




The Sector Specialists

# .net threading

Submitted by: Grzegorz Jachimko  
on: November 2014  
Version: 1.0

Confidential



# Agenda

---

Asynchronous processing in .net

Async model in .net

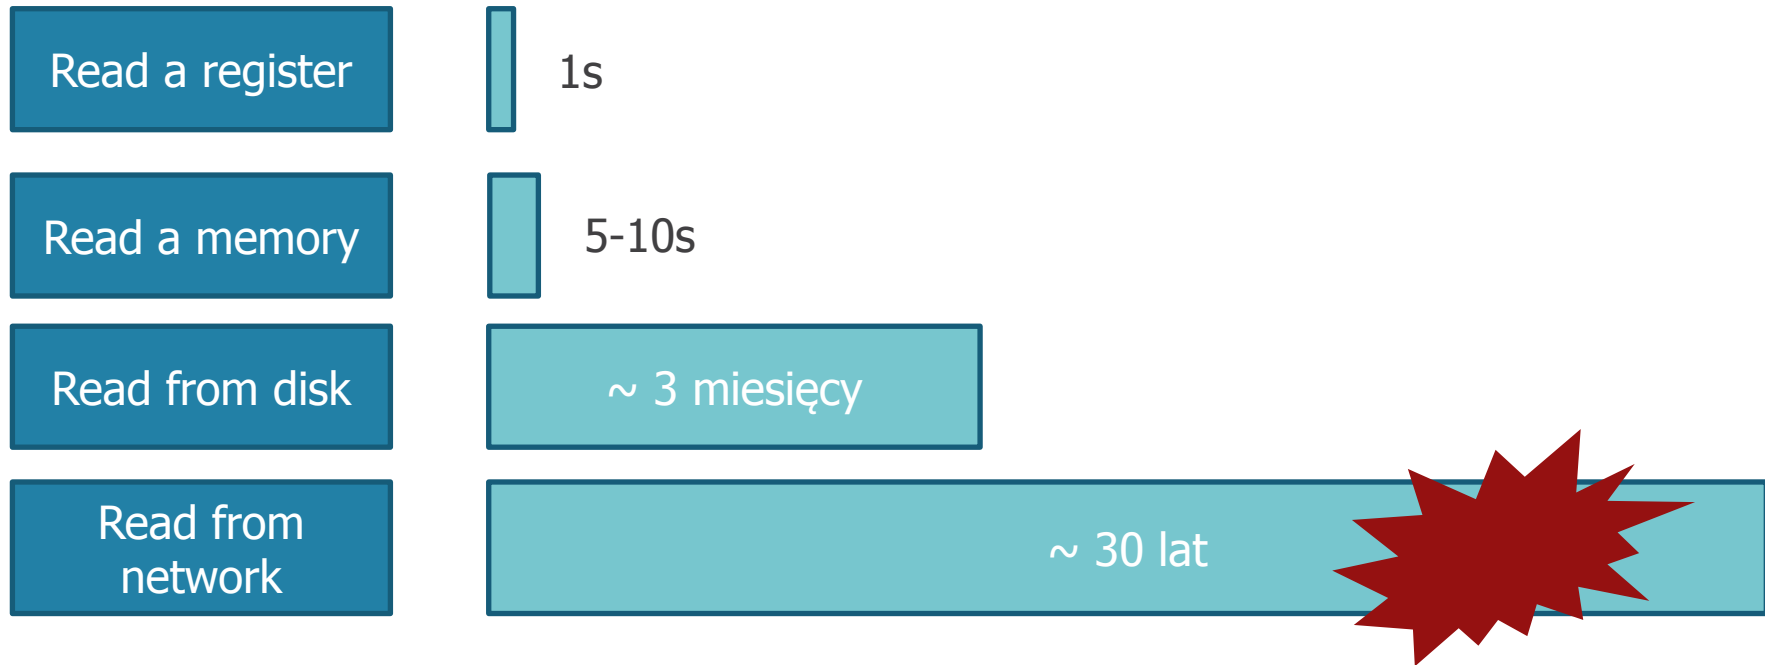
I/O operations

Thread, ThreadPool, Task, async/await

Synchronization

# Some basics

- Processor → 1 cycle = 1 second



# Asynchronous processing in .net

- Windows supports 1+ thread / process to enable multitasking
- But ... thread is a cost
  - ▶ Kernel object (thread's properties, registers – a few KB)
  - ▶ Stack (1MB user mode, ~12/24 KB kernel mode)
  - ▶ DLL attach and detach
  - ▶ Time slots - ~20 ms per thread
- Switching between threads:
  - ▶ Enter kernel mode
  - ▶ Store register values
  - ▶ Select next runnable thread
  - ▶ Load register values
  - ▶ Exit kernel mode

# Asynchronous processing in .net

## DEMO

How many threads can you run in 32 bit process?



# Asynchronous processing in .net

## Threads – cons vs pros

- Context switching is a waste of time
  - ▶ ... but required to keep the system robust and responsive ☹️
- So?
  - ▶ Avoid threads whenever possible,
  - ▶ Use threads to ensure scalability and responsiveness
- Ideally – 1 process (1 thread) per processor
- Reality:



Process Explorer - Sysinternals: www.sysinternals.com [REDMOND\gjachim]

File Options View Process Find Users Help

Process	PID	CPU	Threads	Private Bytes	Working Set
winamp.exe	6704	0.38	31	18 920 K	33 36
wmialm.exe	2364		27	59 444 K	76 65
common32.exe	732		5	1 988 K	2 80
procexp64.exe	8468	1.15	8	34 580 K	48 31
devenv.exe	21340		62	243 780 K	286 58
ThreadsDemo.vshost.exe	13960		7	16 384 K	17 13
mstsc.exe	8728		14	34 672 K	46 76
explore.exe	8948		10	8 292 K	23 62
iexplore.exe	1600		17	22 528 K	37 25
communicator.exe	9052	0.38	42	44 684 K	32 83
cmd.exe	6660		1	2 724 K	3 53
firefox.exe	11692		22	90 152 K	113 95
plugin-container.exe	17356		10	16 092 K	19 10
bacmonitor.exe	6472		3	4 240 K	1 28
MSBuild.exe	8592		36	100 468 K	82 84
cmd.exe	19464		1	2 340 K	3 55
MSTestShim.exe	13120	24.83	17	35 832 K	37 83
VSTestHost.exe	9024		17	22 804 K	18 76
MSBuild.exe					5 86

CPU Usage: 59.97% Commit Charge: 47.04% Processes: 115 Physical Usage: 71.76%

# Async model in .net

## Definition

- Asynchronous programming model is a **KEY** to create **scalable** and **non-blocking** applications
  - ▶ Threads are not blocked while waiting for the operation to complete
    - Blocked thread still consumes resources,
    - Blocked thread causes more threads to be created,
    - More threads = more hassle when running GC

# Async model in .net

## Types of asynchronous operations

- **CPU bound operations** – a set of instructions to be executed by the processor only
- **I/O bound** – hardware (disk, network card, dvd player) has to fetch some data for us – thread has nothing to do!





# Async model in .net

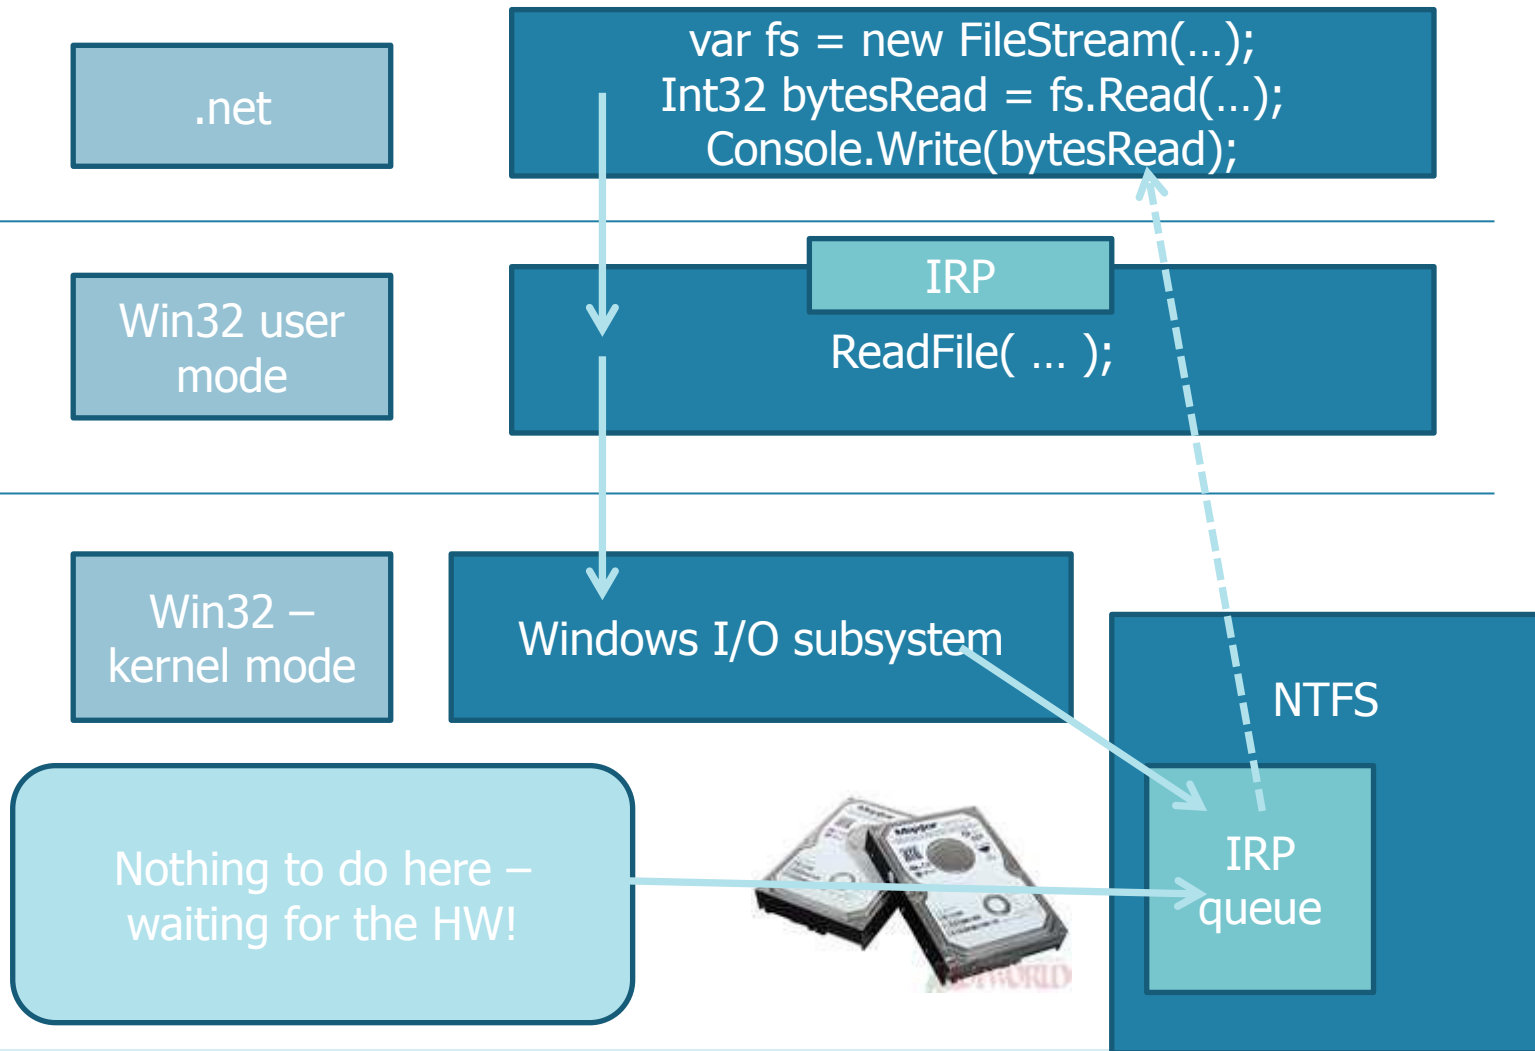
## DEMO

- Present the difference between IO and CPU bound operations



# I/O operations

## How does windows handle synchronous I/O?



# I/O operations

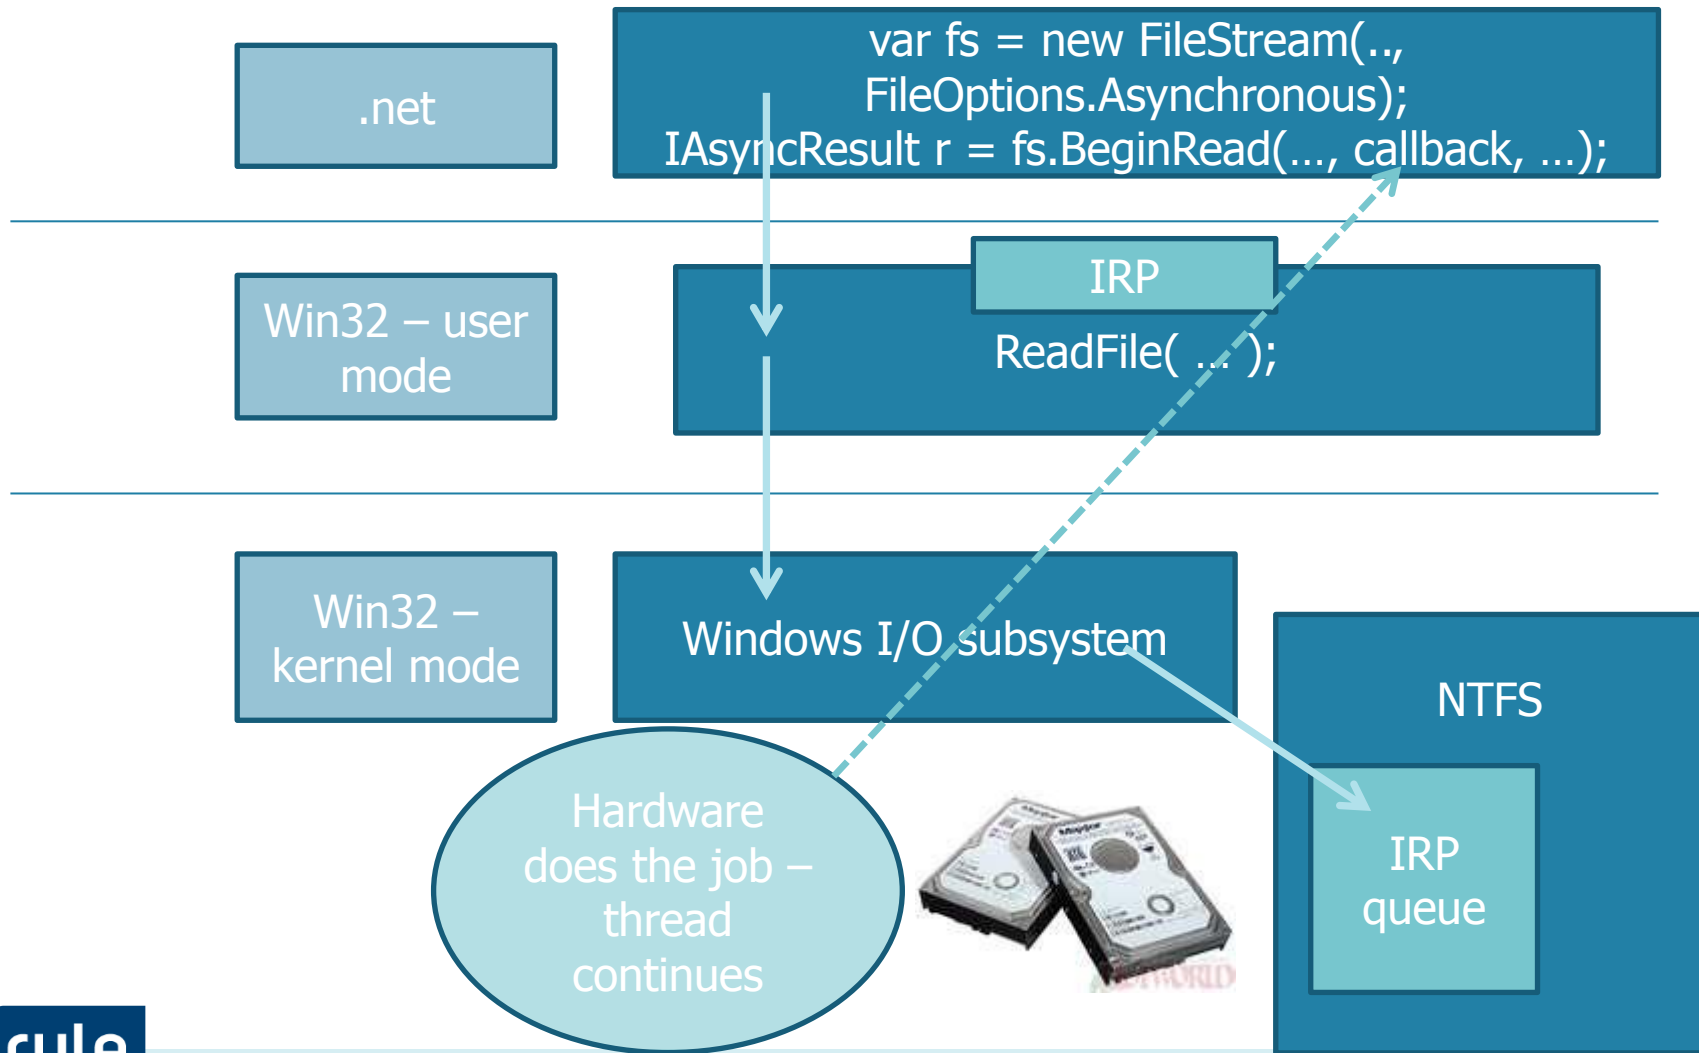
## Why this is important?

- 2 main goals:
  - ▶ Do not block a thread while performing I/O operation
  - ▶ Keep the processor busy with least threads possible
- Windows I/O completion port
- I/O operations done in async:
  - ▶ Not blocking calling thread
  - ▶ Completed IRP get to I/O CP (CLR)
  - ▶ Each completed operation needs a thread assigned



# I/O operations

## How does windows handle asynchronous I/O?



# I/O operations

## DEMO

- How to read a file using APM ?



# Thread, ThreadPool, Task, async/await

## What we have here?

### ● Threads

- ▶ Full control over creation and configuration,
- ▶ Expensive (only when you call Start)
- ▶ Use them – if your operation is long running (>500 ms) or does blocking I/O
- ▶ `new Thread(ThreadStart threadFunction)`



# Thread, ThreadPool, Task, async/await

## What we have here?

### ● ThreadPool

- ▶ CLR controls lifetime
- ▶ Do not change thread's properties (priority, name, etc.)
- ▶ Suggested for short, CPU bound operations (<500 ms)
- ▶ `ThreadPool.QueueUserWorkItem(WaitCallback callback, object state)`



# Thread, ThreadPool, Task, async/await

## What we have here?

### ● TaskPool

- ▶ CLR controls lifetime
- ▶ Abstraction on top of ThreadPool, same rules apply
- ▶ Optimizations for multi-core environments
- ▶ Supports reading operation's value!
- ▶ `Task.Factory.StartNew<TResult>(Func<TResult> func)`






# Thread, ThreadPool, Task, async/await

## What we have here?

### ● Async/await

- ▶ Syntactic sugar on top of Tasks
- ▶ Provides out of the box synchronization


```
▶ private Task<int> Foo() { ... }  
private async Task<int> UseFoo() {  
    int result = await Foo();  
    return result;  
}
```



# Synchronization

## Why?


- Avoid deadlocks
  - ▶ Occurs when 2 threads try to lock their own resources
- Avoid out-of-order execution
  - ▶ Occurs when code gets optimized and is invoked in multithreaded environment
- Avoid data corruption
  - ▶ Occurs when one thread is writing data, and other is trying to read it



# Synchronization

## What options do we have

- Interlocked methods
  - ▶ Add, Exchange, CompareExchange, ... - no blocking, ensure operation is atomic
- Critical section
  - ▶ Great for in-process communication, acquiring is lightning fast\* (`lock (_obj) { ... }`)
- WaitHandles (Manual|Auto ResetEvent, Semaphore, Mutex)
  - ▶ Huge performance impact, but great control capabilities
- UI synchronization
  - ▶ UI applications have STA model, all access operations need to be synchronized
    - SynchronizationContext



# Critical section

## Incrementation problem

```
class Program {  
    private static int index;  
  
    static void Main(string[] args) {  
        index = 0;  
  
        ThreadPool.QueueUserWorkItem(IncrementLoop);  
        ThreadPool.QueueUserWorkItem(IncrementLoop);  
        ThreadPool.QueueUserWorkItem(IncrementLoop);  
  
        Console.ReadLine();  
    }  
  
    private static void IncrementLoop(Object state) {  
        for (int i = 0; i < 20; i++)  
            Console.WriteLine(index++);  
    }  
}
```

# Critical section

---

- Piece of code that accesses a shared resource
- Cannot be concurrently accessed by more than one thread
- Requests for access from threads are queued
- Requires synchronization mechanism to ensure exclusive access – e.g. semaphore

# Synchronization mechanisms

## Monitor vs Semaphore

- About Semaphore:

- ▶ Limits the number of threads that can access a resource or pool of resources concurrently by using `WaitOne` and `Release` methods.
- ▶ There is no guaranteed order, such as FIFO or LIFO, in which blocked threads enter the semaphore.
- ▶ A thread can enter semaphore multiple times by calling `WaitOne`.
- ▶ Does not enforce thread identity on calls to `WaitOne` or `Release` – it is programmer's responsibility to ensure that threads do not release semaphore too many times.
- ▶ There are two types of semaphores:
  - Local semaphore – exists within process
  - Named system semaphore – visible through the OS, can be used to synchronize the activity of processes

- About Monitor:

- ▶ Simple synchronization mechanism.
- ▶ Marks a statement block as a critical section.
- ▶ Ensures that one thread does not enter a critical section while another thread is in critical section.
- ▶ Queues threads that try to enter locked code, they will wait until critical section is released.

# Synchronization mechanisms

## Monitor vs Semaphore

### Semaphore

```
private static Semaphore pool;
private static int sum = 0;

private static void Main (int value) {
    pool = new Semaphore(0, 3);
    for (int i = 0; i < 25; i++)
        ThreadPool.QueueUserWorkItem(Sum, i);
}

private void Sum(object state) {
    int value = (int)state;

    pool.WaitOne();
    sum += value ;
    pool.Release();
}
```

### Monitor

```
private static object syncLock;
private static int sum = 0;

private static void Main (int value) {
    for (int i = 0; i < 25; i++)
        ThreadPool.QueueUserWorkItem(Sum, i);
}

private void Sum(object state) {
    int value = (int)state;

    Monitor.Enter(syncLock);
    sum += value ;
    Monitor.Exit(syncLock);
}
```

# Synchronization mechanisms

## Lock & Interlocked

- About `lock`:

- ▶ Is a nice looking wrapper for `Monitor`, calling `Enter` at the start and `Exit` at the end of block of code.
- ▶ Lock can be acquired on a variable of any reference type.

- About `Interlocked`:

- ▶ Provides atomic operations for variables that are shared by multiple threads.
- ▶ Protects against errors that can occur when scheduler switches context while a thread is updating a shared variable, or when multiple threads are executing concurrently on separate processors.
- ▶ Methods of this class do not throw exceptions.



# Synchronization mechanisms

## Lock & Interlocked

### Lock

```
class A {  
    private static readonly object syncLock =  
        new object();  
    private int a = 0;  
  
    public int Add (int value) {  
        lock(syncLock) {  
            a += value;  
  
            return a;  
        }  
    }  
}
```

### Interlocked

```
class A {  
    private int a = 0;  
  
    public int Add (int value) {  
        return Interlocked.Add(ref a, value);  
    }  
}
```

# Synchronization mechanisms

## Mutex

- Synchronization primitive.
- Can be used for interprocess synchronization.
- Grants exclusive access to shared resource to only one thread.
- One thread can request ownership multiple times using `WaitOne` without blocking its execution. However must call `ReleaseMutex` the same number of times to release ownership.
- Enforce thread identity, so a mutex can be released only by the thread that acquired it, as opposed to Semaphore.
- Like semaphore, there are two types of mutexes:
  - ▶ Local (unnamed) mutex – exists only within a process
  - ▶ Named system mutex – visible through the OS and can be used to synchronize the activities of processes

# Synchronization mechanisms

## Mutex

```
private static Mutex mutex;
private static int sum = 0;

private static void Main (int value) {
    pool = new Mutex();

    for (int i = 0; i < 25; i++)
        ThreadPool.QueueUserWorkItem(Sum, i);
}

private void Sum(object state) {
    int value = (int)state;

    mutex.WaitOne();
    sum += value ;
    mutex.ReleaseMutex();
}
```

# Synchronization mechanisms

## AutoResetEvent & Task

- About `AutoResetEvent`:

- ▶ Notifies a waiting thread that an event has occurred, which allows threads to communicate with each other by signaling.
- ▶ Thread waits for signal by calling non-blocking loop – `WaitOne`.
- ▶ Calling `Set` signals waiting thread.
- ▶ Can be used with static `WaitAll` and `WaitAny` methods.

- Synchronization using `Task`:

- ▶ `WaitAll` is similar to `AutoResetEvent.WaitAll` – it is non-blocking loop waiting for all tasks to complete execution.
- ▶ `WaitAny` is similar to `AutoResetEvent.WaitAny` – it is non-blocking loop waiting for any of the tasks to complete execution.
- ▶ `WhenAll` creates a new task that will complete when all of the provided tasks have completed.
- ▶ `WhenAny` creates a new task that will complete when any of the provided tasks have completed.

# Multithreading issues

## Deadlock

### Coffman conditions

- Mutual exclusion – two or more resources must not be shareable, and only one process can use the resource at the time.
- Hold and wait or resource holding – process is currently holding at least one resource and requesting additional resources which are being held by other processes.
- No preemption – resources must be released by holding process voluntarily, they must not be deallocated by operating system.
- Circular wait – process must be waiting for a resource which is being held by another process, which in turn is waiting for first process to release the resource.

# Multithreading issues

## Deadlock

```
private void Transfer(Account accountA, Account accountB, double amount)
{
    lock (accountA)
    {
        lock (accountB)
        {
            if (accountA.Balance < amount)
                throw new Exception("Insufficient funds.");

            accountA.Balance -= amount;
            accountA.Balance += amount;
        }
    }
}
```

# Multithreading issues

## Deadlock - solution

```
private void Transfer(Account accountA, Account accountB, double amount)
{
    if (accountA.GetHashCode() > accountB.GetHashCode())
        lock (accountA)
            lock (accountB)
            {
                // logic
            }
    else
        lock (accountB)
            lock (accountA)
            {
                // logic
            }
}
```

# Multithreading issues

## Race condition

```
class A {  
    int result = 0;  
    void Work1() { result = 1; }  
    void Work2() { result = 2; }  
    void Work3() { result = 3; }  
  
    static void Main(string[] args) {  
        A a = new A();  
        Thread worker1 = new Thread(a.Work1);  
        Thread worker2 = new Thread(a.Work2);  
        Thread worker3 = new Thread(a.Work3);  
        worker1.Start();  
        worker2.Start();  
        worker3.Start();  
        Console.WriteLine(a.result);  
        Console.Read();  
    }  
}
```

Result of the code might be 1, 2, 3 or even 0.



# Good practices

---

- Prefer using methods of Interlocked class for simple state changes, instead of using lock keyword.
- Use producer/consumer model.
- Do not use types as lock objects.
- Use caution when locking on instances, e.g. lock(this).
- Do ensure that a thread that has entered a monitor always leaves that monitor even if an exception occurs while thread is in the monitor, lock statement provides this behavior automatically, employing a finally block to ensure that Monitor.Exit() is called.

# References

---

- <http://msdn.microsoft.com/en-us/magazine/cc164040.aspx>
- <http://stackoverflow.com/questions/154551/volatile-vs-interlocked-vs-lock>
- <http://www.pzielinski.com/?p=1692>
- <http://www.c-sharpcorner.com/UploadFile/1d42da/race-conditions-in-threading-C-Sharp/>
- <https://www.wintellect.com/blogs/jeffreyr>
- <http://www.softwareinteractions.com/blog/2009/12/8/intrprocess-message-queue.html>

**Q & A**